

Assignment_2_EH

October 17, 2018

1 Assignment 2

1.0.1 MACS 3000, Dr. Evans

1.0.2 Ellen Hsieh

```
In [1]: import numpy as np
        from sklearn import linear_model
        import pandas as pd
        import matplotlib.pyplot as plt
        from sklearn.metrics import mean_squared_error, r2_score
```

1.0.3 1. Imputing age and gender

(a) **My proposed strategy to impute age and gender into BestIncome.txt:** First, we build two linear regression models respect to age and gender based on total income and weight in dataset SurvIncome.txt respectively. Then, we use those two models to predict the value of age and gender in BestIncome.txt. However, there is no feature called total income in BestIncome.txt, therefore, we can create the value of total income by adding labor and capital income in BestIncome.txt. Finally, we use the total income and weight in BestIncome.txt to calculate the value of age and gender for BestIncome.txt based on those two regression models. The equations are as following:

$$Age_i = \beta_0 + \beta_1 Total_income_i + \beta_2 Weight_i + \epsilon_i$$

$$Gender_i = \beta_0 + \beta_1 Total_income_i + \beta_2 Weight_i + \epsilon_i$$

```
In [2]: # open and rename the column names of the imported files
        best = pd.read_csv('BestIncome.txt', sep=',', header=None)
        survey = pd.read_csv('SurvIncome.txt', sep=',', header=None)
        best = best.rename(index=int, columns={0:'lab_inc', 1:'cap_inc', 2:'hgt', 3:'wgt'})
        survey = survey.rename(index=int, columns={0:'tot_inc', 1:'wgt', 2:'age', 3:'gender'})
```

```
In [3]: best.describe()
```

```
Out [3]:
```

	lab_inc	cap_inc	hgt	wgt
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	57052.925133	9985.798563	65.014021	150.006011
std	8036.544363	2010.123691	1.999692	9.973001

min	22917.607900	1495.191896	58.176154	114.510700
25%	51624.339880	8611.756679	63.652971	143.341979
50%	56968.709935	9969.840117	65.003557	149.947641
75%	62408.232277	11339.905773	66.356915	156.724586
max	90059.898537	19882.320069	72.802277	185.408280

In [4]: `survey.describe()`

```
Out [4]:
```

	tot_inc	wgt	age	gender
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	64871.210860	149.542181	44.839320	0.500000
std	9542.444214	22.028883	5.939185	0.50025
min	31816.281649	99.662468	25.741333	0.000000
25%	58349.862384	130.179235	41.025231	0.000000
50%	65281.271149	149.758434	44.955981	0.500000
75%	71749.038000	170.147337	48.817644	1.000000
max	92556.135462	196.503274	66.534646	1.000000

(b) Here is where I'll use my proposed method from part (a) to impute variables.

```
In [5]: # create the total income value by adding labor income and capital income
tot_inc_best = best.lab_inc + best.cap_inc
x_best = np.column_stack((tot_inc_best, best.wgt))
```

```
In [6]: # with the common features: tot_inc and wgt
# age ~ tot_inc and wgt
x = np.column_stack((survey.tot_inc, survey.wgt))
y_age = survey.age.values
y_gender = survey.gender.values

# simple linear regression: age ~ tot_inc + wgt
regr_age = linear_model.LinearRegression()
regr_age.fit(x, y_age)

# logistic regression: gender ~ tot_inc + wgt
regr_gender = linear_model.LogisticRegression()
regr_gender.fit(x, y_gender)
```

```
Out [6]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False)
```

```
In [7]: # compute the predicted age and gender based on the information in best
pred_age = regr_age.predict(x_best)
pred_gender = regr_gender.predict(x_best)
```

```
In [8]: # simple append the predictions onto the best dataset
imputed_best = np.column_stack((best, pred_age, pred_gender))
```

(c) Here is where I'll report the descriptive statistics for my new imputed variables.

```
In [9]: # convert the numpy array into a pandas dataframe
new_best = pd.DataFrame({'lab_inc':imputed_best[:,0], 'cap_inc':imputed_best[:,1],
                        'hgt':imputed_best[:,2], 'wgt':imputed_best[:,3],
                        'age':imputed_best[:,4], 'gender':imputed_best[:,5]})
```

```
In [10]: new_best.describe()
```

```
Out[10]:
```

	lab_inc	cap_inc	hgt	wgt	age
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	57052.925133	9985.798563	65.014021	150.006011	44.890828
std	8036.544363	2010.123691	1.999692	9.973001	0.219150
min	22917.607900	1495.191896	58.176154	114.510700	43.976495
25%	51624.339880	8611.756679	63.652971	143.341979	44.743776
50%	56968.709935	9969.840117	65.003557	149.947641	44.886944
75%	62408.232277	11339.905773	66.356915	156.724586	45.038991
max	90059.898537	19882.320069	72.802277	185.408280	45.703819

	gender
count	10000.000000
mean	0.471700
std	0.499223
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

(d) Correlation matrix for the now six variables

```
In [11]: # Correlation matrix form
corr = new_best.corr()
corr.style.background_gradient()
```

```
Out[11]: <pandas.io.formats.style.Styler at 0x1a162981d0>
```

```
In [12]: # Correction Matrix Plot
def corr_plot(df):

    names = df.columns
    N = len(names)


    correlations = df.corr()
    fig = plt.figure()
    ax = fig.add_subplot(111)
    cax = ax.matshow(correlations, vmin=-1, vmax=1)
    fig.colorbar(cax)
    ticks = np.arange(0,N,1)
```

```

ax.set_xticks(ticks)
ax.set_yticks(ticks)
ax.set_xticklabels(names)
ax.set_yticklabels(names)
plt.show()

```

In [13]: `corr_plot(new_best)`



1.0.4 2. Stationarity and data drift

(a) Estimate by OLS and report coefficients

```

In [14]: # open and rename the column names of the imported files
inc_intel = pd.read_csv('IncomeIntel.txt', sep=',', header=None)
inc_intel = inc_intel.rename(index=int, columns={0:'grad_year', 1:'gre_qnt', 2:'salary'})

In [15]: # Define Outcome and Independent Variables
outcome = 'salary_p4'
features = ['gre_qnt']

X, y = inc_intel[features], inc_intel[outcome]

In [16]: # run ols regression and report the result
import statsmodels.api as sm

X_vars = X[['gre_qnt']]

```

```

X_vars = sm.add_constant(X_vars, prepend=False)
X_vars.head()

m = sm.OLS(y, X_vars)

res = m.fit()
print(res.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:                  salary_p4      R-squared:                0.263
Model:                            OLS      Adj. R-squared:            0.262
Method:                 Least Squares      F-statistic:                356.3
Date:                  Tue, 16 Oct 2018      Prob (F-statistic):        3.43e-68
Time:                  23:33:29      Log-Likelihood:            -10673.
No. Observations:                1000      AIC:                      2.135e+04
Df Residuals:                    998      BIC:                      2.136e+04
Df Model:                        1
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
gre_qnt	-25.7632	1.365	-18.875	0.000	-28.442	-23.085
const	8.954e+04	878.764	101.895	0.000	8.78e+04	9.13e+04

```

=====
Omnibus:                        9.118      Durbin-Watson:                1.424
Prob(Omnibus):                  0.010      Jarque-Bera (JB):              9.100
Skew:                          0.230      Prob(JB):                      0.0106
Kurtosis:                      3.077      Cond. No.                      1.71e+03
=====

```

Warnings:


- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.71e+03. This might indicate that there are strong multicollinearity or other numerical problems.

(b) Create a scatterplot of GRE score and graduation year.

```

In [17]: # make a scatterplot for gre score and graduation year
grad_year = inc_intel['grad_year']
gre_qnt = inc_intel['gre_qnt']
inc_intel.plot(x='grad_year', y='gre_qnt', kind='scatter')
plt.show()


```



The problem here is that there is a obvious jump down on this scatterplot due to the different scales of GRE score. Before 2011, the score scale for GRE quatative was 200-800 and after that it became 130-170. Therefore, if we use this dataset to build a regression model to predict the salary, it would distort the result of linear regression. In order to solve this problem, we can convert all the GRE scores to the latter scale(130-170), so that it could become consistent. And then, we could get a more precise regression model based on the modified dataset. (Note: The conversion used is not based on percentile so the value might not be precise enough).


```
In [18]: # convert gre_qnt score
converted_gre_qnt = inc_intel.gre_qnt[:]
for i in range(len(inc_intel)):
    if inc_intel.grad_year[i] <= 2010.0:
        converted_gre_qnt[i] = 130+ ((inc_intel.gre_qnt[i]-200)*40/600)
    else:
        pass

In [19]: # a new scatterplot for converted gre score
grad_year = inc_intel['grad_year']
gre_qnt = converted_gre_qnt
inc_intel.plot(x='grad_year', y='gre_qnt', kind='scatter')
plt.show()
```



(c) Create a scatterplot of income and graduation year

```
In [20]: # create a scatterplot for income and graduation year
grad_year = inc_intel['grad_year']
salary_p4 = inc_intel['salary_p4']
inc_intel.plot(x='grad_year', y='salary_p4', kind='scatter')
plt.show()
```



The problem here is that the income is increasing over years. This is because we didn't take the inflation rate into consideration. To solve this problem, we can convert all the incomes into the present value by divided them by the inflation rate regarding to different year.


```
In [21]: # calculate the average income by year
avg_inc_by_year = inc_intel['salary_p4'].groupby(inc_intel['grad_year']).mean().values

# calculate the average growth rate
avg_growth_rate = ((avg_inc_by_year[1:] - avg_inc_by_year[:-1]) / avg_inc_by_year[:-1])

In [22]: # calcualte the present value of all the incomes
new_salary_p4 = salary_p4[:]

for i in range(len(salary_p4)):
    if inc_intel.grad_year[i] == 2001.0:
        new_salary_p4[i] = inc_intel.salary_p4[i]
    else:
        new_salary_p4[i] = inc_intel.salary_p4[i] / ((1 + avg_growth_rate)**(inc_intel.grad_year[i] - 2001.0))

In [23]: # a new scatterplot for new income value
grad_year = inc_intel['grad_year']
salary_p4 = new_salary_p4
inc_intel.plot(x='grad_year', y='salary_p4', kind='scatter')
plt.show()
```

(d) Re-estimate coefficients with updated variables.

```
In [24]: # add the modified features to inc_intel
inc_intel['converted_gre_qnt'] = converted_gre_qnt
inc_intel['new_salary_p4'] = new_salary_p4

# define the outcome and features for regression model
outcome = 'new_salary_p4'
features = ['converted_gre_qnt']

X, y = inc_intel[features], inc_intel[outcome]

# run the regression for new variables
X_vars = X[['converted_gre_qnt']]
X_vars = sm.add_constant(X_vars, prepend=False)
X_vars.head()

m = sm.OLS(y, X_vars)

res = m.fit()
print(res.summary())
```

OLS Regression Results

=====

```

Dep. Variable:          new_salary_p4    R-squared:                0.000
Model:                  OLS              Adj. R-squared:           -0.001
Method:                 Least Squares    F-statistic:              0.05257
Date:                   Tue, 16 Oct 2018  Prob (F-statistic):      0.819
Time:                   23:33:30         Log-Likelihood:           -10291.
No. Observations:       1000            AIC:                     2.059e+04
Df Residuals:           998            BIC:                     2.060e+04
Df Model:                1
Covariance Type:        nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
converted_gre_qnt    -9.9681      43.474     -0.229      0.819     -95.279      75.343
const                6.304e+04    7083.395      8.900      0.000     4.91e+04     7.69e+04
=====
Omnibus:                 0.757    Durbin-Watson:                2.026
Prob(Omnibus):            0.685    Jarque-Bera (JB):            0.668
Skew:                     0.059    Prob(JB):                    0.716
Kurtosis:                 3.048    Cond. No.                     5.11e+03
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.11e+03. This might indicate that there are strong multicollinearity or other numerical problems.

The result of OLS regression for modified GRE score and salary shows that the new estimated coefficients became -9.9681 which is less negative than the former one. From this fact, we can know that GRE score now has less negative impact on salary. Before the conversion of GRE score and salary, the larger range of GRE score and abnormally distribution made GRE score have larger negative effect on salary. Therefore, after modified the dataset, the correlation between those variables become more accurate and the negative impact from GRE score on salary become less significant, which supports my hypothesis for the result of the modified regression model.

1.0.5 3. Assessment of Kossinets and Watts.

See attached PDF.